

```
In [45]: import pandas as pd
import numpy as np
```

```
In [46]: data=pd.read_csv("/home/placement/Downloads/TelecomCustomerChurn.csv")
```

```
In [47]: data.head()
```

```
Out[47]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtect
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	

5 rows × 21 columns

In [61]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   float64
20  Churn                  7043 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

```
In [49]: data.isna().sum()
```

```
Out[49]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport    0
StreamingTV    0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

```
In [50]: data.dtypes
```

```
Out[50]: customerID      object
gender                  object
SeniorCitizen          int64
Partner                object
Dependents              object
tenure                  int64
PhoneService            object
MultipleLines           object
InternetService         object
OnlineSecurity          object
OnlineBackup            object
DeviceProtection        object
TechSupport             object
StreamingTV             object
StreamingMovies         object
Contract                object
PaperlessBilling        object
PaymentMethod           object
MonthlyCharges          float64
TotalCharges            object
Churn                   object
dtype: object
```

```
In [51]: data['TotalCharges'] = pd.to_numeric(data['TotalCharges'],errors='coerce')
#data['TotalCharges'] = data['TotalCharges'].replace('',np.nan).astype(float).value
```

In [52]: data.dtypes

```
Out[52]: customerID      object
gender      object
SeniorCitizen  int64
Partner      object
Dependents   object
tenure       int64
PhoneService  object
MultipleLines object
InternetService object
OnlineSecurity object
OnlineBackup  object
DeviceProtection object
TechSupport   object
StreamingTV   object
StreamingMovies object
Contract      object
PaperlessBilling object
PaymentMethod object
MonthlyCharges float64
TotalCharges  float64
Churn         object
dtype: object
```

```
In [53]: #from sklearn.model_selection import GridSearchCV #GridSearchCV is for parameter tuning
#from sklearn.ensemble import RandomForestClassifier
#cls=RandomForestClassifier()
#n_estimators=[25,50,75,100,125,150,175,200] #number of decision trees in the forest, default = 100
#criterion=['gini','entropy'] #criteria for choosing nodes default = 'gini'
#max_depth=[3,5,10] #maximum number of nodes in a tree default = None (it will go till all possible nodes)
#parameters={'n_estimators': n_estimators,'criterion':criterion,'max_depth':max_depth} #this will undergo 8
#RFC_cls = GridSearchCV(cls, parameters)
#RFC_cls.fit(X_train,y_train)
```

In [54]: databackup=data.copy()

In [55]: data['TotalCharges']=data['TotalCharges'].fillna(data['TotalCharges'].median())

```
In [56]: x=data.drop(['customerID','Churn'],axis=1)
y=data['Churn']
```

```
In [57]: x=pd.get_dummies(x)
```

```
In [58]: x.head()
```

```
Out[58]:
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_Female	gender_Male	Partner_No	Partner_Yes	Dependents_No	Dependents_Yes
0	0	1	29.85	29.85	1	0	0	1	1	0
1	0	34	56.95	1889.50	0	1	1	0	1	0
2	0	2	53.85	108.15	0	1	1	0	1	0
3	0	45	42.30	1840.75	0	1	1	0	1	0
4	0	2	70.70	151.65	1	0	1	0	1	0

5 rows × 45 columns

In [63]: `list(data)`

Out[63]: ['customerID',
'gender',
'SeniorCitizen',
'Partner',
'Dependents',
'tenure',
'PhoneService',
'MultipleLines',
'InternetService',
'OnlineSecurity',
'OnlineBackup',
'DeviceProtection',
'TechSupport',
'StreamingTV',
'StreamingMovies',
'Contract',
'PaperlessBilling',
'PaymentMethod',
'MonthlyCharges',
'TotalCharges',
'Churn']

In [59]: `from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)`

```
In [60]: from sklearn.model_selection import GridSearchCV #GridSearchCV is for parameter tuning
from sklearn.ensemble import RandomForestClassifier
cls=RandomForestClassifier()
n_estimators=[25,50,75,100,125,150,175,200] #number of decision trees in the forest, default = 100
criterion=['gini','entropy'] #criteria for choosing nodes default = 'gini'
max_depth=[3,5,10] #maximum number of nodes in a tree default = None (it will go till all possible nodes)
parameters={'n_estimators': n_estimators,'criterion':criterion,'max_depth':max_depth} #this will undergo 8*2
RFC_cls = GridSearchCV(cls, parameters)
RFC_cls.fit(x_train,y_train)
```

```
Out[60]: GridSearchCV(estimator=RandomForestClassifier(),
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_depth': [3, 5, 10],
                                   'n_estimators': [25, 50, 75, 100, 125, 150, 175, 200]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [64]: RFC_cls.best_params_
```

```
Out[64]: {'criterion': 'gini', 'max_depth': 10, 'n_estimators': 50}
```

```
In [66]: cls=RandomForestClassifier(n_estimators=50,criterion='entropy',max_depth=10)
```

```
In [67]: cls.fit(x_train,y_train)
```

```
Out[67]: RandomForestClassifier(criterion='entropy', max_depth=10, n_estimators=50)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [68]: rfy_pred=cls.predict(x_test)
```

```
In [69]: rfy_pred
```

```
Out[69]: array(['Yes', 'No', 'No', ..., 'Yes', 'No', 'No'], dtype=object)
```



```
In [70]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test, rfy_pred)
```

```
Out[70]: array([[1538, 159],
               [ 300, 328]])
```

```
In [71]: from sklearn.metrics import accuracy_score
accuracy_score(y_test, rfy_pred)
```

```
Out[71]: 0.8025806451612904
```

```
In [72]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression()
classifier.fit(x_train, y_train)
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[72]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [73]: y_pred = classifier.predict(x_test)
```

```
In [74]: from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

```
Out[74]: 0.8120430107526881
```

In []: